

# BDA - Assignment 8

Anonymous

5/4/2020

```
library(tidyverse)
library(aaltobda)
library(rstan)
# stan settings:
source('stan_utility.R') # diagnosis of rhats
options(mc.cores = parallel::detectCores()) #for local computer
rstan_options(auto_write = TRUE) # autosave Stan
# bay settings:
library(loo) #pred. error of MCMC log likelihood
library(gridExtra)
library(bayesplot) #plots of posterior draws (mcmc_hist etc)
library(shinystan) # model paramteres & MCMC simulations
bayesplot_theme_set() #default
SEED <- 48927 # random seed for reproducability
data("factory")
```

In this assignment we will do a model assessment using the leave-one-out cross-validation (LOO-CV) for the factory data from the aaltobda package. In the factory data, the quality of 6 machines are given.

## Q1 - Fitting of the models

Given we have 6 machines, our three stan models are:

```
# display the stanmodels
writeLines(readLines("ex8_separate.stan"))

## // seperate model for factory data
## data {
##   int<lower=0> N; // number of data points
##   int<lower=0> K; // number of groups
##   int<lower=1,upper=K> x[N]; // group indicator
##   vector[N] y; //
## }
## parameters {
##   vector[K] mu; // group means
##   vector<lower=0>[K] sigma; // stds of group
## }
## model {
##   y ~ normal(mu[x], sigma[x]);
## }
```

```

## generated quantities {
##   real ypred;
##   vector[N] log_lik;
##   ypred = normal_rng(mu[6], sigma[6]);
##   for (i in 1:N)
##     log_lik[i] = normal_lpdf(y[i] | mu[x[i]], sigma[x[i]]);
## }

```

```
writeLines(readLines("ex8_pooled.stan"))
```

```

## // pooled model for factory data
## data {
##   int<lower=0> N; // number of data points
##   vector[N] y; //
## }
## parameters {
##   real mu; // prior means
##   real<lower=0> sigma; // prior std
## }
## model {
##   y ~ normal(mu, sigma);
## }
## generated quantities {
##   real ypred;
##   vector[N] log_lik;
##   ypred = normal_rng(mu, sigma);
##   for (i in 1:N)
##     log_lik[i] = normal_lpdf(y[i] | mu, sigma);
## }

```

```
writeLines(readLines("ex8_hierarchical.stan"))
```

```

## // hierarchical model for factory standard
## data {
##   int<lower=0> N; // number of data points
##   int<lower=0> K; // number of groups
##   int<lower=1,upper=K> x[N]; // group indicator
##   vector[N] y; //
## }
## parameters {
##   real mu0; // prior mean
##   real<lower=0> sigma0; // prior std
##   vector[K] mu; // group means
##   real<lower=0> sigma; // common stds
## }
## model {
##   mu0 ~ normal(90, 15); // weakly informative prior
##   sigma0 ~ cauchy(0,4); // weakly informative prior
##   mu ~ normal(mu0, sigma0); // population prior with unknown parameters
##   sigma ~ cauchy(0,4); // weakly informative prior
##   y ~ normal(mu[x], sigma);
## }
## generated quantities {

```

```
##   real ypred;
##   real mu7;
##   vector[N] log_lik;
##   ypred = normal_rng(mu[6], sigma);
##   mu7 = normal_rng(mu0, sigma0);
##   for (i in 1:N)
##     log_lik[i] = normal_lpdf(y[i] | mu[x[i]], sigma);
## }
```

### Fitting of the seperate model:

Assumptions:

- Priors are uniform
- There are 6 machines
- Each machine (j) has unrelated means  $\mu_j$  and standard deviations  $\sigma_j$

```
# creating data
d_separate <-list(N = 30,
                 K = 6,
                 x = rep(1:ncol(factory), nrow(factory)),
                 y = c(t(factory)))

# fitting
fit_separate <- stan(file="ex8_separate.stan",
                    data = d_separate, seed = SEED)
```

### Fitting of the pooled model:

Assumptions:

- Priors are uniform
- The machines equals only one machine, where all the data comes from
- The one machine has one mean  $\mu$  and one standard deviation  $\sigma$

```
# creating data
d_pooled <- list(N = 30,
                y = c(t(factory)))

# fitting
fit_pooled <- stan(file = "ex8_pooled.stan",
                  data = d_pooled, seed = SEED)
```

### Fitting of the hierarchical model:

Assumptions:

- Prior distribution follows:  $\mu \sim \mathcal{N}(\mu_0, \sigma_0)$
- Equal standard deviation for all machines  $\sigma$

```
# creating data
d_hierarchical <- list(N = 30,
                      K = 6,
                      x = rep(1:ncol(factory), nrow(factory)),
                      y = c(t(factory)))

# fitting
fit_hierarchical <- stan(file="ex8_hierarchical.stan",
                        data = d_hierarchical, seed = SEED)
```

```
## Warning: There were 39 divergent transitions after warmup. Increasing adapt_delta above 0.8 may help
## http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
```

```
## Warning: Examine the pairs() plot to diagnose sampling problems
```

```
## Warning: Tail Effective Samples Size (ESS) is too low, indicating posterior variances and tail quantiles
## Running the chains for more iterations may help. See
## http://mc-stan.org/misc/warnings.html#tail-ess
```

## Q2

For model checking and comparison of models, we will do a Pareto smoothed importance-sampling leave-one-out cross-validation (PSIS-LOO) to compute the expected log predictive density (elpd) values and  $\hat{k}$ -values. We use the functions from the loo package for this and for a diagnostic plot, visualizing the  $\hat{k}$ -values. The horizontal lines in the plot helps determining if the  $\hat{k}$ -values are good or not. A value  $<5$  is considered good,  $<0.7$  is ok and  $>0.7$  is bad.

### Separate

```
# seperate
log_lik_separate <- extract_log_lik(fit_separate,
                                   merge_chains = FALSE)
r_eff_separate <- relative_eff(exp(log_lik_separate))
loo_separate <- loo(log_lik_separate,
                   r_eff = r_eff_separate)
```

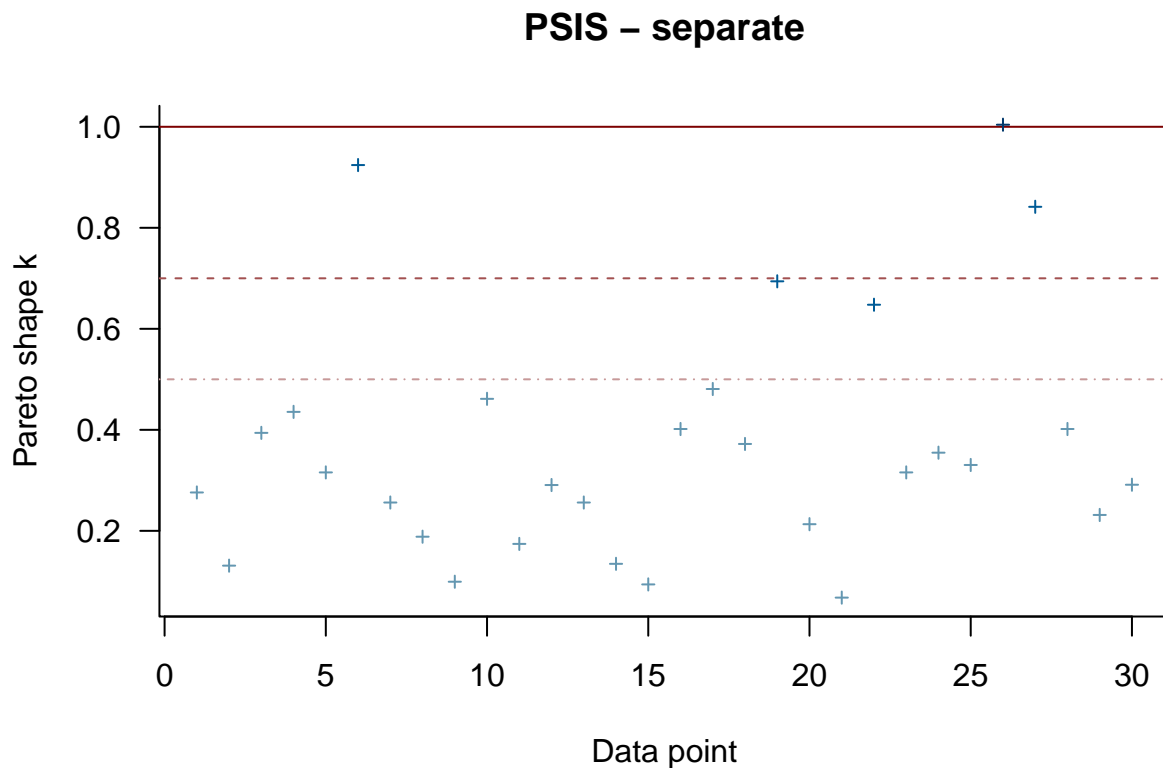
```
## Warning: Some Pareto k diagnostic values are too high. See help('pareto-k-diagnostic') for details.
```

```
print(loo_separate)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate SE
## elpd_loo   -132.6 3.2
## p_loo       9.9 1.1
## looic      265.2 6.3
## -----
## Monte Carlo SE of elpd_loo is NA.
##
```

```
## Pareto k diagnostic values:
##                               Count Pct.   Min. n_eff
## (-Inf, 0.5]   (good)         25   83.3%   442
## (0.5, 0.7]   (ok)           2    6.7%   274
## (0.7, 1]     (bad)           2    6.7%    77
## (1, Inf)     (very bad)      1    3.3%    64
## See help('pareto-k-diagnostic') for details.
```

```
# visualizing
plot(loo_separate, diagnostic = c("k", "n_eff"),
     label_points = FALSE, main = "PSIS - separate")
```



For the separate model we get:

- PSIS-LOO : -132.6
- $\hat{k}$  : Pareto k diagnostic values are slightly high (10% is above 0.7).

## Pooled


```
# pooled
log_lik_pooled <- extract_log_lik(fit_pooled,
                                merge_chains = FALSE)
r_eff_pooled <- relative_eff(exp(log_lik_pooled))
loo_pooled <- loo(log_lik_pooled,
                 r_eff = r_eff_pooled)
```

```
## Warning: Some Pareto k diagnostic values are slightly high. See help('pareto-k-diagnostic') for details.
```

```
print(loo_pooled)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate SE
## elpd_loo   -131.0 4.3
## p_loo       2.1 0.8
## looic      262.1 8.6
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## Pareto k diagnostic values:
##           Count Pct.    Min. n_eff
## (-Inf, 0.5] (good)    29   96.7%   1896
## (0.5, 0.7]  (ok)      1    3.3%    510
## (0.7, 1]    (bad)      0    0.0%    <NA>
## (1, Inf)    (very bad) 0    0.0%    <NA>
##
## All Pareto k estimates are ok (k < 0.7).
## See help('pareto-k-diagnostic') for details.
```

```
# visualizing:
plot(loo_pooled, diagnostic = c("k", "n_eff"),
     label_points = FALSE, main = "PSIS - pooled")
```



For the pooled model we get:

- PSIS-LOO: -131.0
- $\hat{k}$  : All Pareto k estimates are ok ( $k < 0.7$ )

**hierachical**

```
# hierarchial
log_lik_hierarchical <-
  extract_log_lik(fit_hierarchical, merge_chains = FALSE)
r_eff_hierarchical <-
  relative_eff(exp(log_lik_hierarchical))
loo_hierarchical <-
  loo(log_lik_hierarchical, r_eff = r_eff_hierarchical)
```

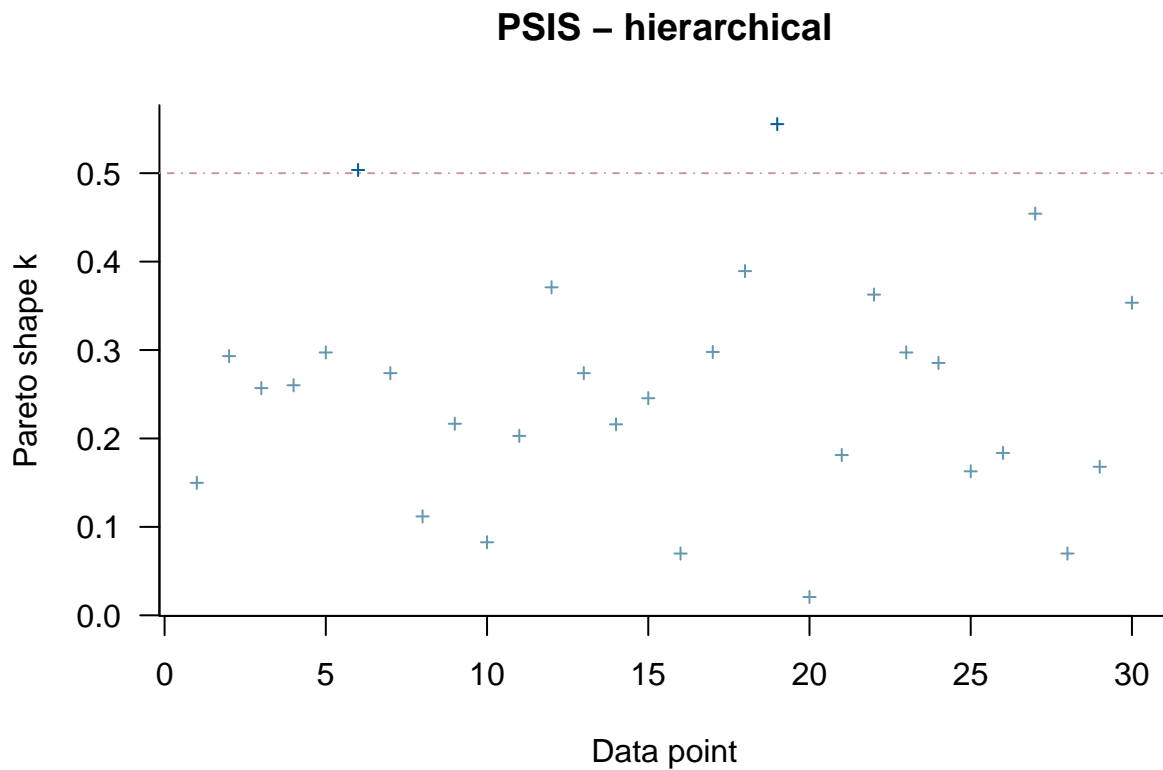
```
## Warning: Some Pareto k diagnostic values are slightly high. See help('pareto-k-diagnostic') for details.
```

```
print(loo_hierarchical)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##      Estimate  SE
```

```
## elpd_loo    -127.2 4.5
## p_loo       5.4 1.5
## looic       254.5 9.0
## -----
## Monte Carlo SE of elpd_loo is 0.1.
##
## Pareto k diagnostic values:
##               Count Pct.    Min. n_eff
## (-Inf, 0.5]  (good)    28   93.3%    682
## (0.5, 0.7]   (ok)      2    6.7%    264
## (0.7, 1]     (bad)      0    0.0%    <NA>
## (1, Inf)     (very bad) 0    0.0%    <NA>
##
## All Pareto k estimates are ok ( $k < 0.7$ ).
## See help('pareto-k-diagnostic') for details.
```

```
# visualizing:
plot(loo_hierarchical, diagnostic = c("k", "n_eff"),
     label_points = FALSE, main = "PSIS - hierarchical")
```



For the hierarchical model we get:

- PSIS-LOO : -127.2
- $\hat{k}$  : All Pareto k estimates are ok ( $k < 0.7$ )



### Q3

The computation of the effective number of parameters  $p_{eff}$  for each of the three models was done in the previous question with the loo function. Printing the results again:

```
print(loo_separate)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate SE
## elpd_loo  -132.6 3.2
## p_loo      9.9 1.1
## looic      265.2 6.3
## -----
## Monte Carlo SE of elpd_loo is NA.
##
## Pareto k diagnostic values:
##           Count Pct.    Min. n_eff
## (-Inf, 0.5] (good)    25    83.3%   442
## (0.5, 0.7] (ok)       2     6.7%   274
## (0.7, 1] (bad)        2     6.7%    77
## (1, Inf) (very bad)   1     3.3%    64
## See help('pareto-k-diagnostic') for details.
```

```
print(loo_pooled)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate SE
## elpd_loo  -131.0 4.3
## p_loo      2.1 0.8
## looic      262.1 8.6
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## Pareto k diagnostic values:
##           Count Pct.    Min. n_eff
## (-Inf, 0.5] (good)    29    96.7%   1896
## (0.5, 0.7] (ok)       1     3.3%    510
## (0.7, 1] (bad)        0     0.0%   <NA>
## (1, Inf) (very bad)   0     0.0%   <NA>
##
## All Pareto k estimates are ok (k < 0.7).
## See help('pareto-k-diagnostic') for details.
```

```
print(loo_hierarchical)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
```

```
##           Estimate SE
## elpd_loo   -127.2 4.5
## p_loo       5.4 1.5
## looic       254.5 9.0
## -----
## Monte Carlo SE of elpd_loo is 0.1.
##
## Pareto k diagnostic values:
##           Count Pct.   Min. n_eff
## (-Inf, 0.5] (good)    28   93.3%   682
## (0.5, 0.7] (ok)       2    6.7%   264
## (0.7, 1] (bad)        0    0.0%   <NA>
## (1, Inf) (very bad)  0    0.0%   <NA>
##
## All Pareto k estimates are ok (k < 0.7).
## See help('pareto-k-diagnostic') for details.
```

## Q4

In Q2 we answered this question with computations and plots. Summarizing the results from Q2:

In the first plot (separated) we see three observations with pareto  $\hat{k} > 0.7$ , which indicates that the model are not reliable. However the two last two diagnostic plots only shows  $\hat{k} < 0.7$ , which means these models are reliable.

## Q5

To compare the models to each other we can use the compare function;

```
compare(loo_separate, loo_pooled)
```

```
## elpd_diff      se
##      1.6      4.0
```

```
compare(loo_separate, loo_hierarchical)
```

```
## elpd_diff      se
##      5.4      3.1
```

```
compare(loo_pooled, loo_hierarchical)
```

```
## elpd_diff      se
##      3.8      1.7
```

We do see a difference between all the models. To summarize the results obtained in Q1-Q4:

### Separate model

- PSIS-LOO : -132.6, the highest of the three models (the smaller the better)
- $\hat{k}$  : 90% good or ok. 10% bad

### **Pooled model**

- PSIS-LOO : -131.
- $\hat{k}$  100 % good or ok

### **Hierarchical model**

- PSIS-LOO : -127.2, smallest of the three models (and hereby best)
- $\hat{k}$  : 100 % good or ok

Overall the most reliable model is the Hierarchical.

### **References:**

Based on code examples from:  
[https://github.com/avehtari/BDA\\_R\\_demos](https://github.com/avehtari/BDA_R_demos)